Timeseries Event Channel

Get insight into your energy usage within hours via AgileDX

2024

CGI

Summary

Datahub

- Business Context
- System Context

Background

- Timeline
- Legislation change
- NextGen meter
- Datahub characteristics
- Stakeholder survey
- Focus Area Technology Study

Technology Study

- Approach
- Patterns
- Requirements
- Technology Analysis
- Design

Proof of Concept

Advise on proof of concepts to be done in phase 2 to determine the feasibility of the chosen architecture and its components.

Summary

Summary

The decree 767/2021 “Valtioneuvoston asetus sähköntoimitusten selvityksestä ja mittauksesta” (Government Decree on the Settlement and Measurement of Electricity Supplies) introduces a directive to make measurement data available within 6 hours per 1/1/2026.

A preliminary analysis of what this could mean for involved parties and the Datahub has been performed and the verification on the feasibility of a high performant event channel has started.

Next step is to closely involve the market to identify the business needs, align on interpretation, align on technicalities including specifications and working models and produce acceptable solutions for affected functionalities.

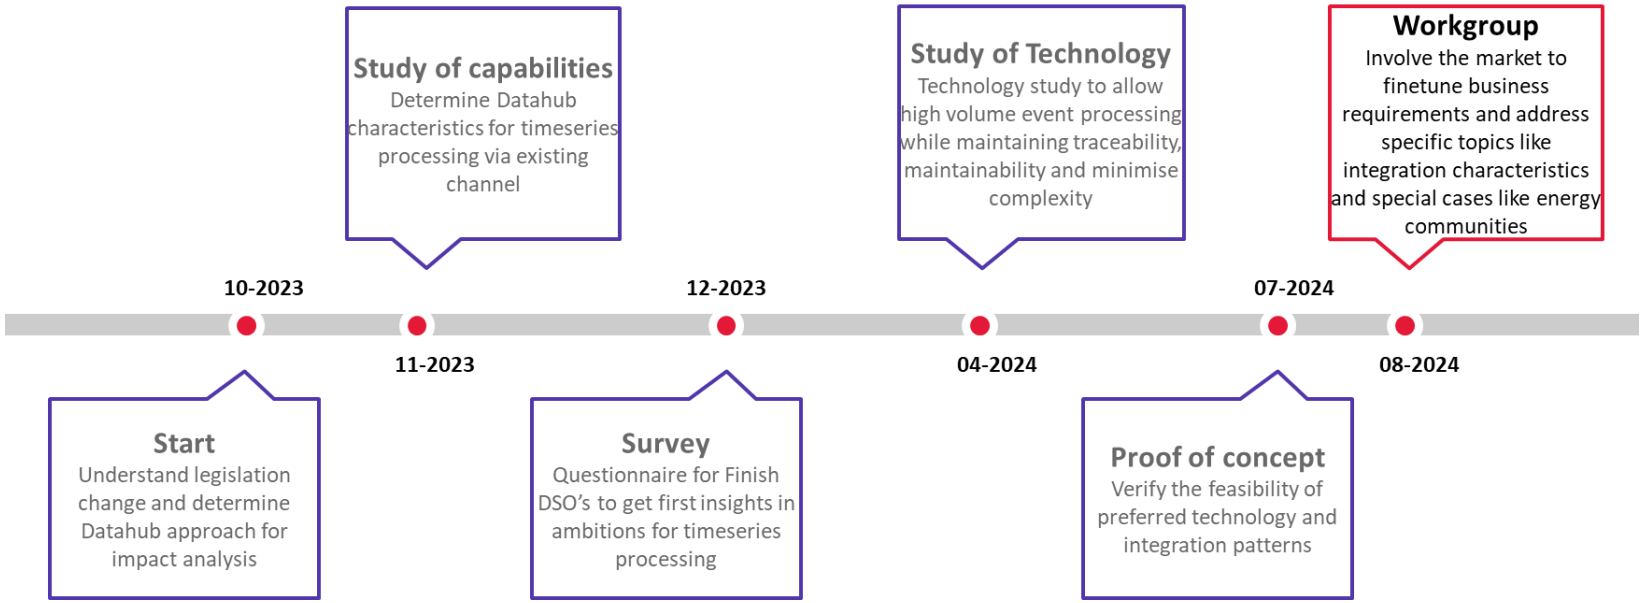

Introduction Datahub

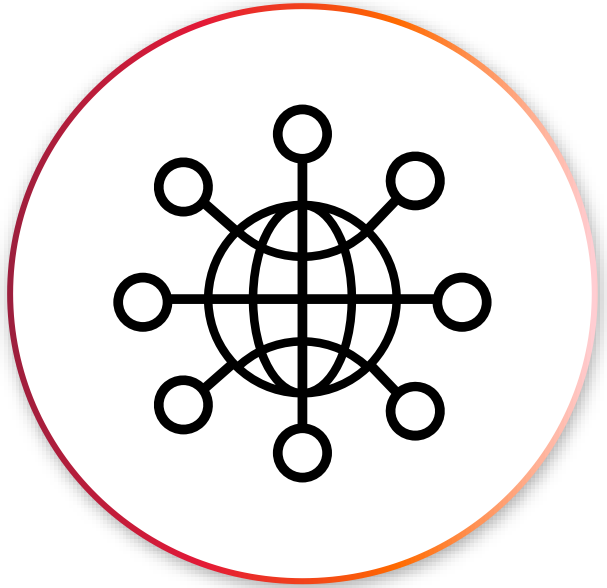

Business Context

Datahub Introduction

Amended in 2013, the Electricity Market Act gave Fingrid Oyj (Fingrid) responsibility for developing the information exchange required for electricity trade and imbalance settlement; the Fingrid Datahub.

Since its implementation, electricity retail market information exchange in Finland is based on a centralised model, in which information exchange between electricity market parties takes place via the Datahub and information is saved in a centralised data storage.

The implementation of this model will trigger a transition from the existing batch processing-based asynchronous information exchange model to real-time synchronous information exchange.

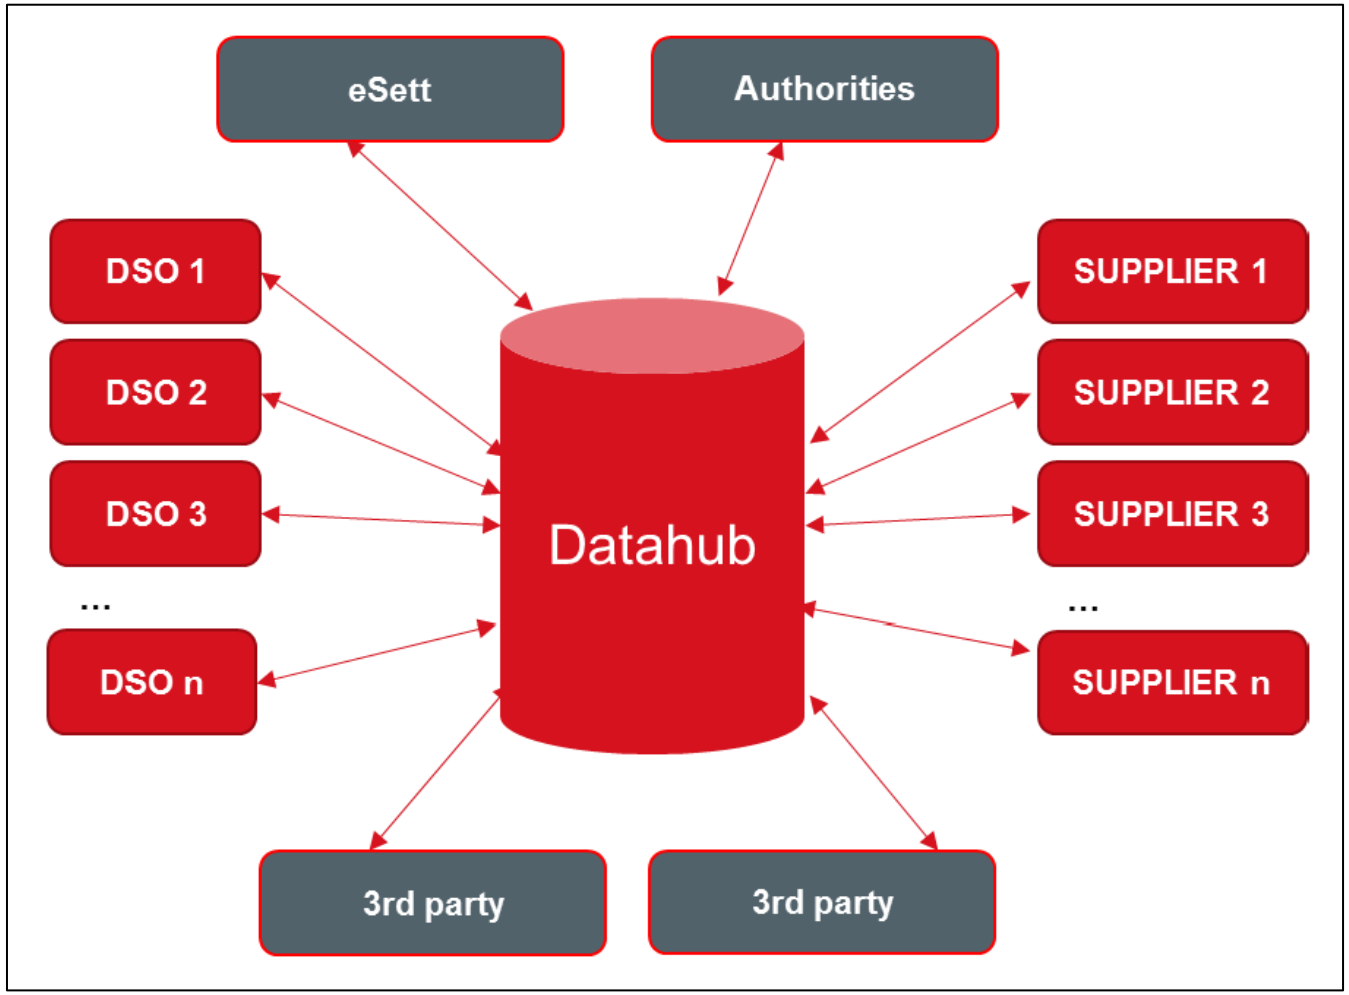

System Context

Introduction Datahub

The block in the middle represents AgileDX (CMS) with its modules.

The surrounding grey/blue boxes are the primary interfaces exposed to market participants:

- Market Party Web Interface – maps on the Web-GUI of CMS, used by all market parties including the market operator.
- Operator User Interface – maps on the Web-GUI of CMS
- Customer Access Interface – maps on the Customer Online Access of CMS
- Data Interface – maps on the B2B interaction of CMS
- B2B User Interface – maps on the B2B interaction of CMS

The purple boxes identify the type of users interacting with the system.

The grey box at the bottom identifies the business processes and other interfaces supported by the solution.

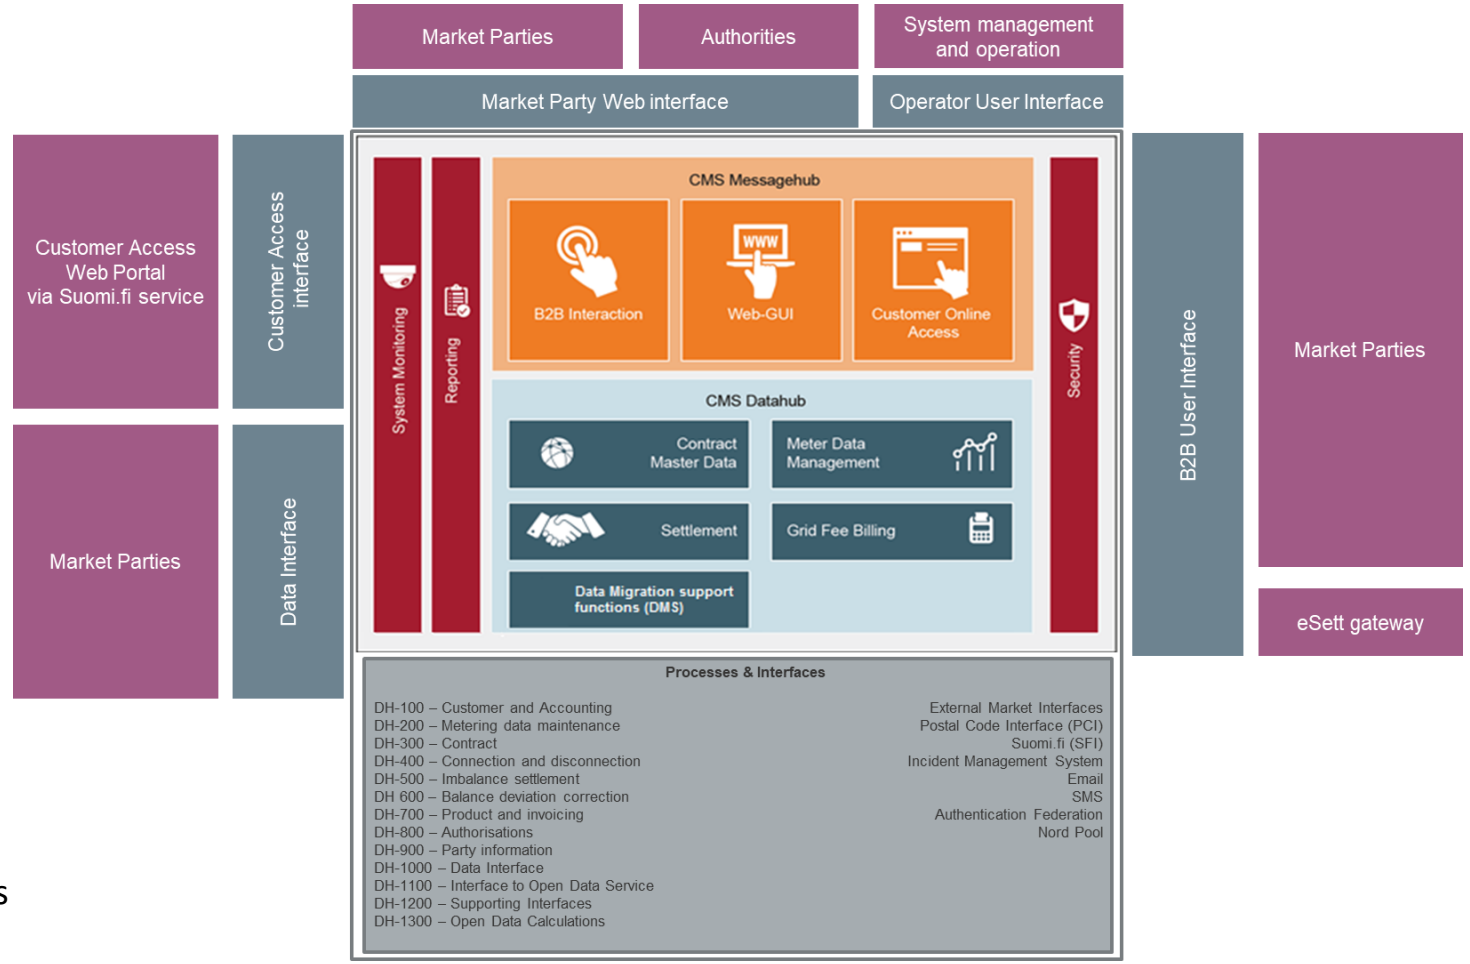

Background

Timeline

Background

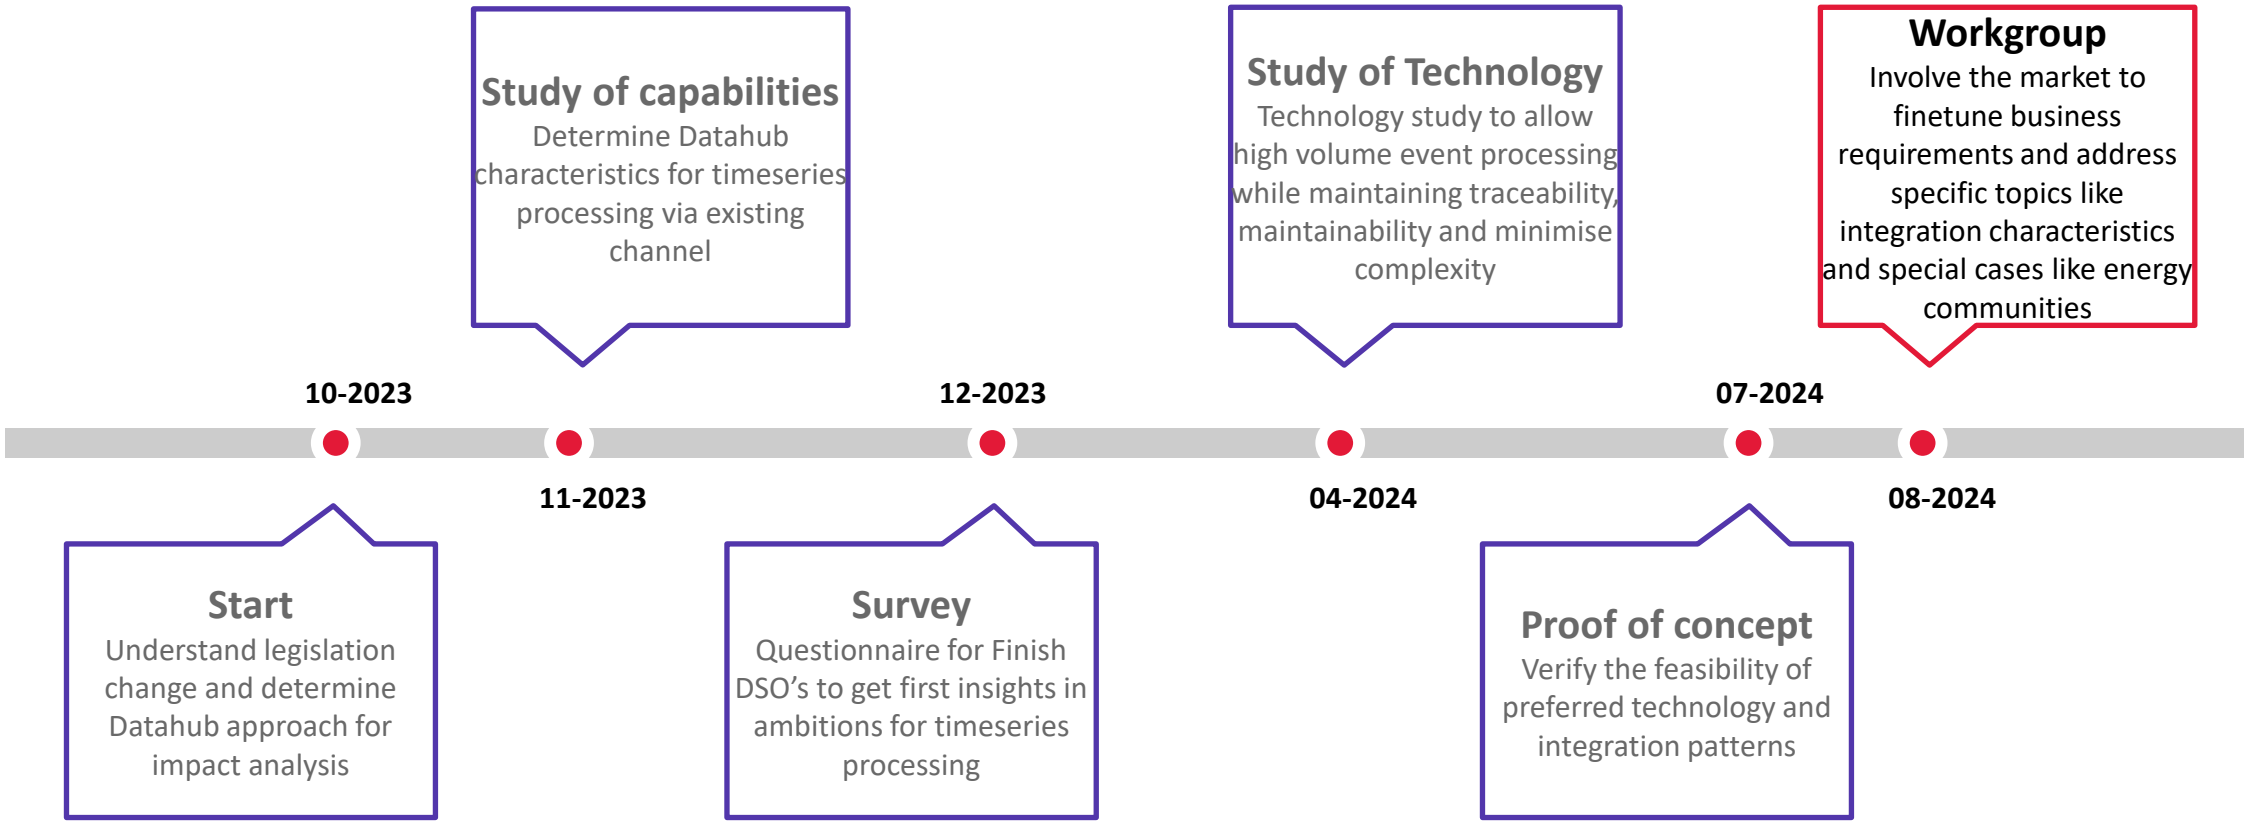

Legislation

Background

The decree 767/2021 “Valtioneuvoston asetus sähköntoimitusten selvityksestä ja mittauksesta” (Government Decree on the Settlement and Measurement of Electricity Supplies) chapter 6, § 5 states

“network operator's information system processing metering data shall collect the registered measurement data from the new remote metering equipment into the metering data reading system at least every six hours”.

According to the Finnish regulator the purpose of the requirement is **to provide measurement data to the end consumer at least within six hours**. As **datahub** can be appointed as point of delivery for measurement data the decree requirement also applies to the datahub system.

The requirement is limited to the **next generation** of smart meters and comes into force **1.1.2026**.

NextGen Meters

Background

GEN I

First Generation meters only provide APIs to request for measurement data and do not have the capability to push measurement data.

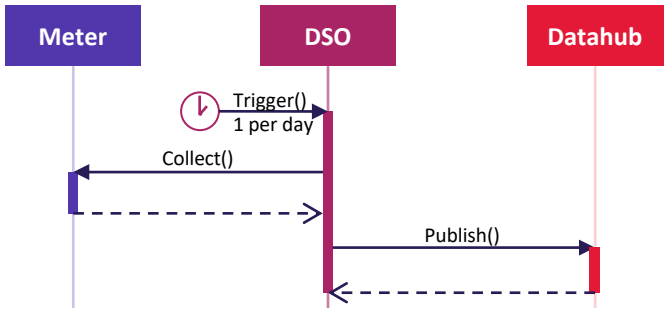

A batch driven approach is more suitable by collecting measurement for a period, validate and then publish for other consumers

GEN II

Second Generation meters have the capability to push the available measurement data as it becomes available.

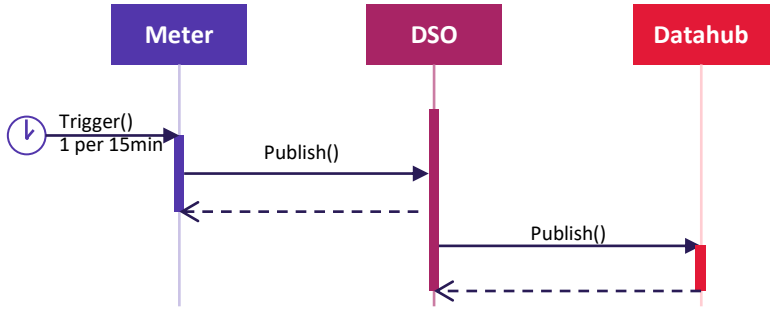

This allows for a more event driven approach to the collection, validation a publication of the available measurement data

Datahub characteristics – Study of Capabilities

Background

The existing channel for receiving timeseries data can be categorized as a batch interface. Designed for current Fingrid situation with additional validations optimized for current integration characteristics.

As a result the current interface is expected to be able to process up to ~250k – 500k messages per day.

Currently Datahub processes around 75k messages a day which in total contains around 500 mio timeseries

Fingrid Survey

Background

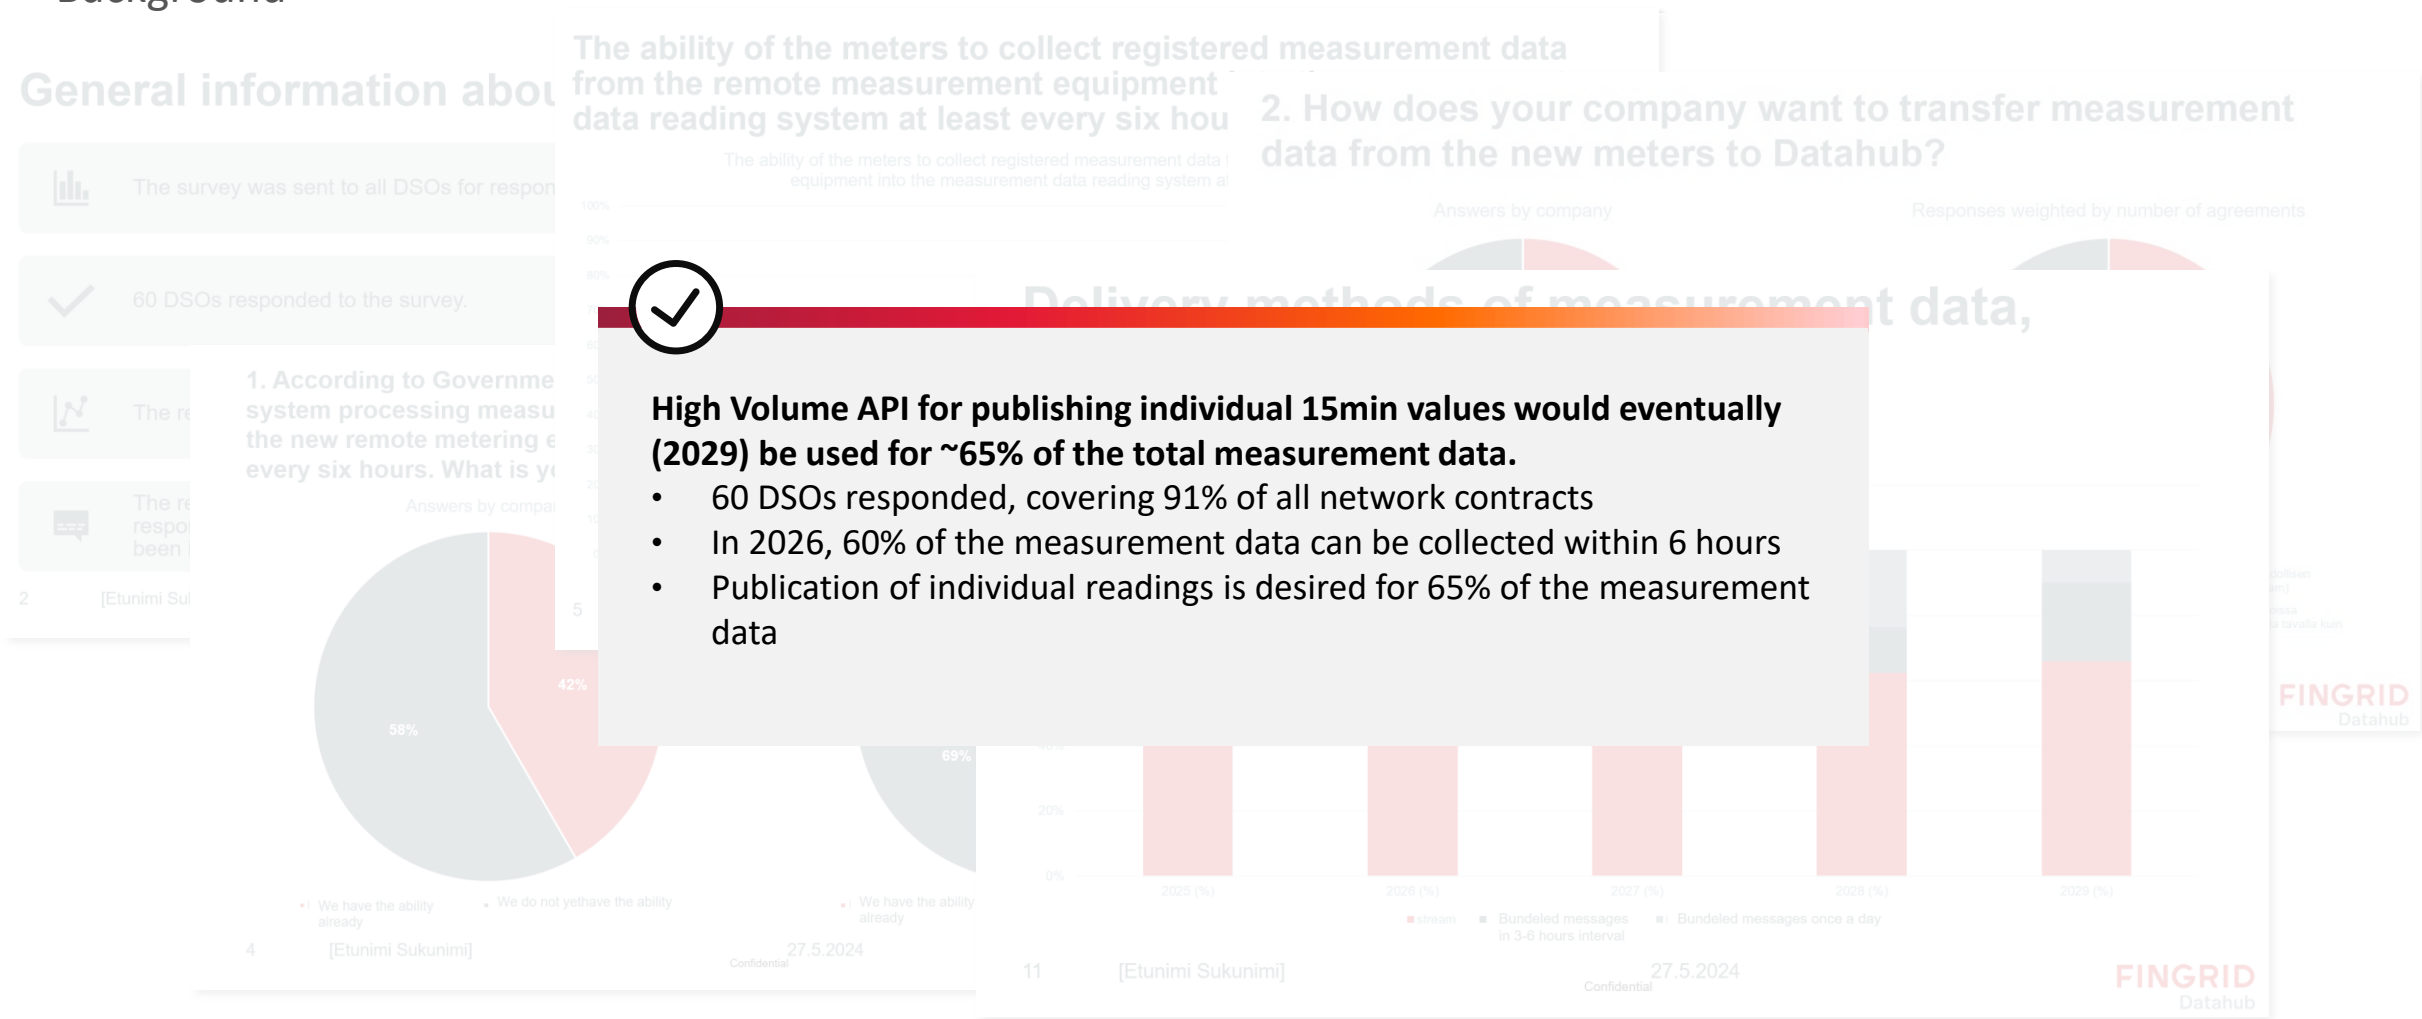

High Volume API for publishing individual 15min values would eventually (2029) be used for ~65% of the total measurement data.

- 60 DSOs responded, covering 91% of all network contracts
- In 2026, 60% of the measurement data can be collected within 6 hours
- Publication of individual readings is desired for 65% of the measurement data

Focus Area Technology Study

Background

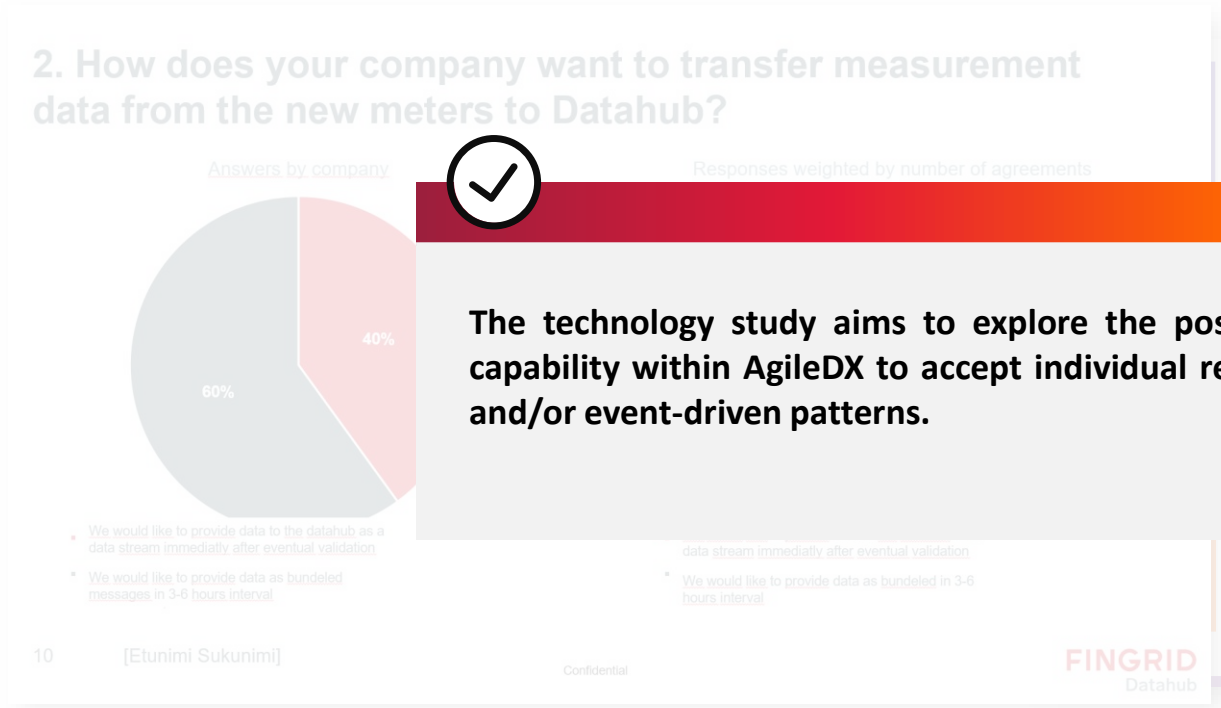

The technology study aims to explore the possibility of integrating the capability within AgileDX to accept individual readings through streaming and/or event-driven patterns.

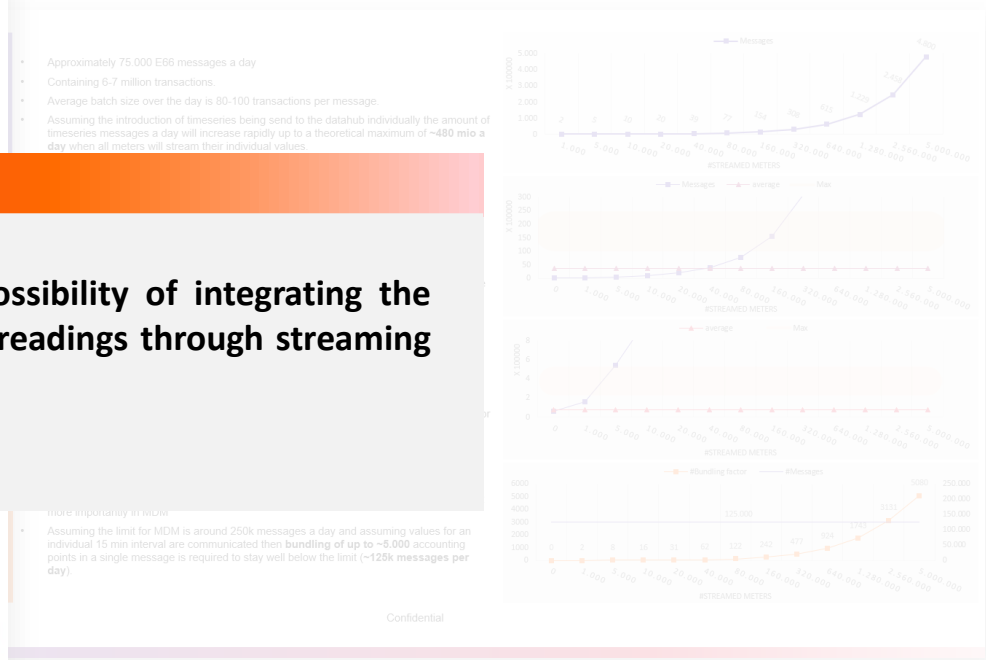

Technology Study

Approach

Technology Study

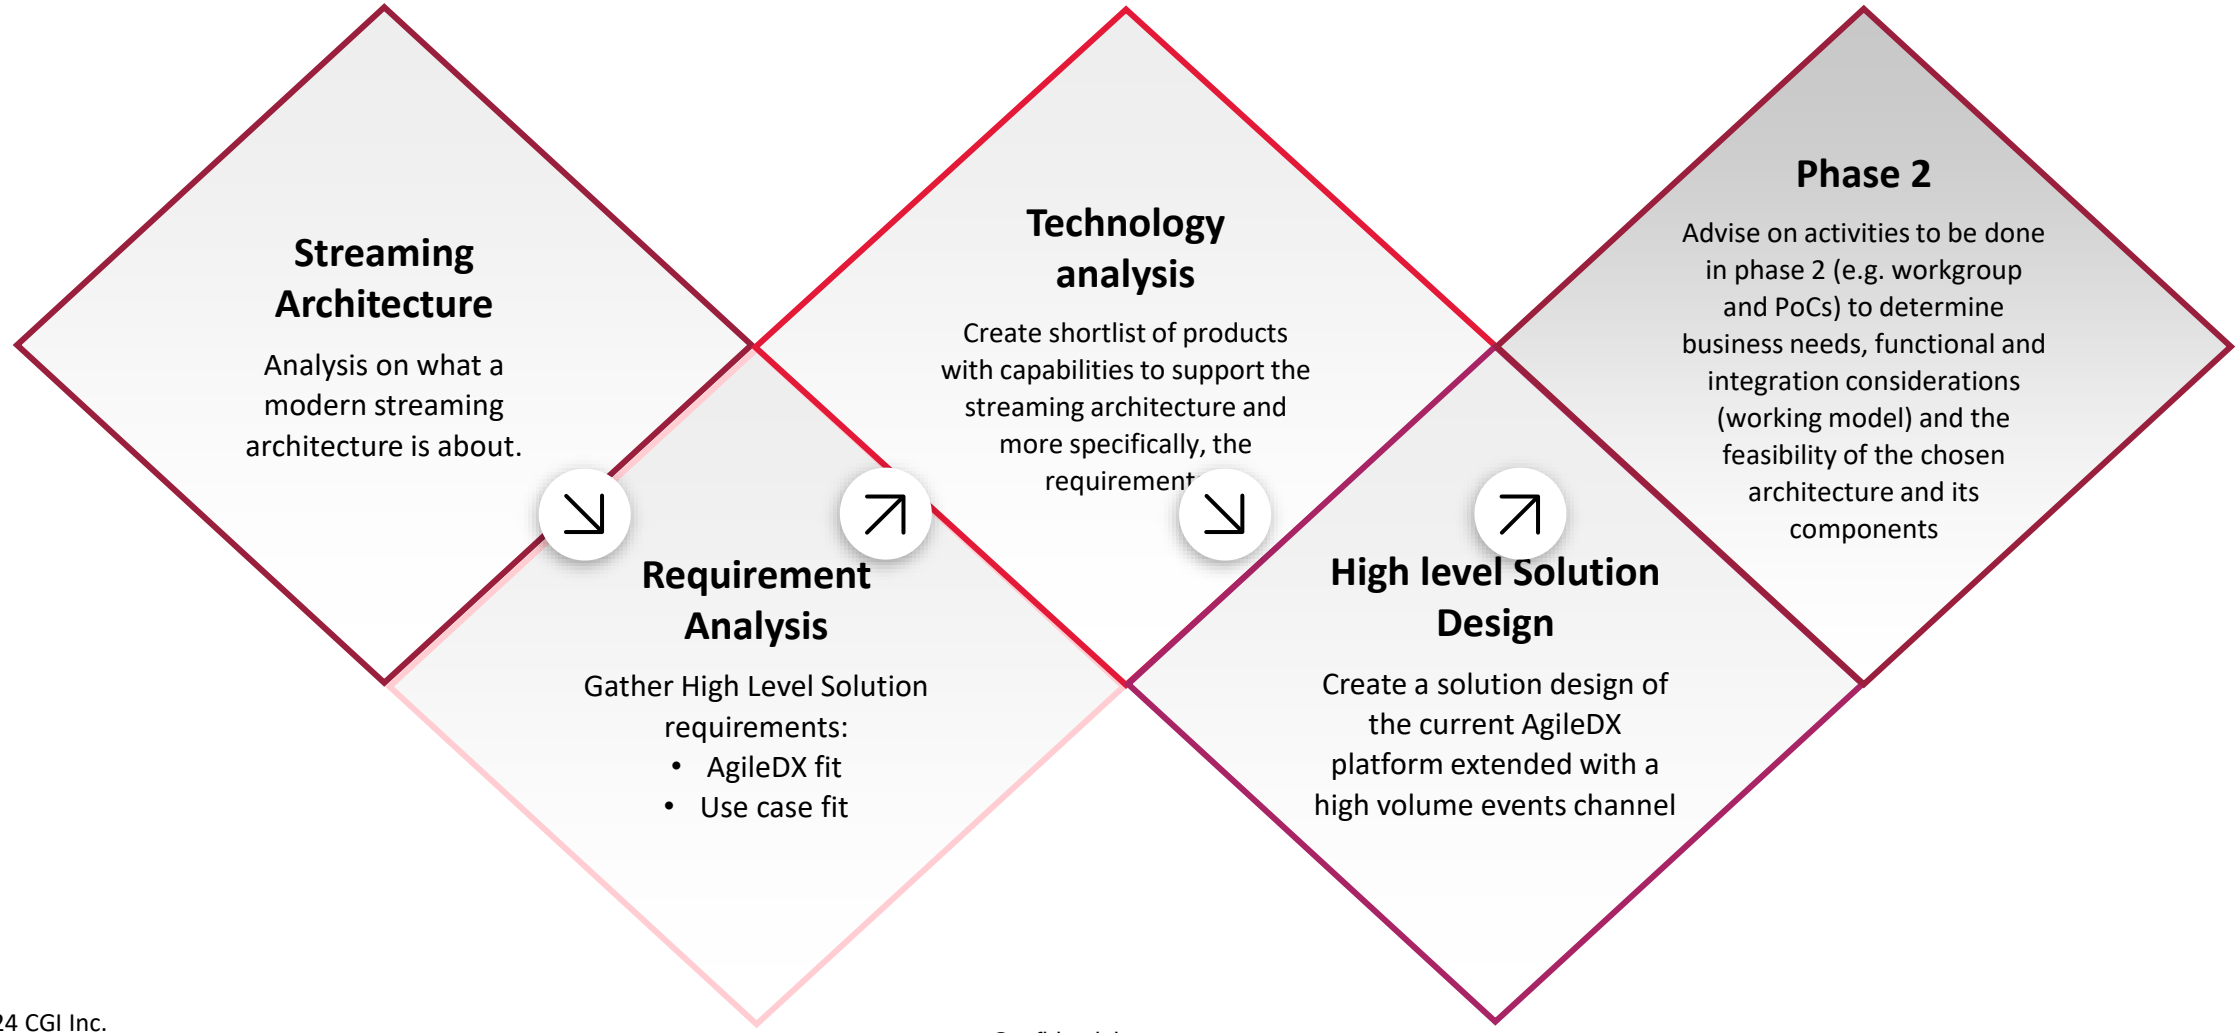

What is streaming data

Technology Study

Streaming data is data that is emitted at high volume in a continuous, incremental manner with the goal of low-latency processing.

Characteristics

Chronologically significant
Individual elements in a data stream contain time stamps

Continuously flowing
A data stream has no beginning or end

Unique
Repeat transmission of a data stream is challenging because of time sensitivity

Nonhomogeneous
Some sources may stream data in multiple formats that are in structured formats

Imperfect
Temporary errors at the source may result in damaged or missing elements in the streamed data

Source: *©Amazon - https://aws.amazon.com/what-is/streaming-data/*

Modern Data Streaming Architecture

Technology Study

Stream Sources	Stream Ingestion	Stream Storage	Stream Processing	Stream Destination
Your source of streaming data includes data sources like sensors, social media, IoT devices, log files generated by using your web and mobile applications, mobile devices that generates semi-structured and unstructured data as continuous streams at high velocity.	The stream ingestion layer is responsible for **ingesting data into the stream storage layer**. It provides the ability to collect data from tens of thousands of data sources and ingest in **near real-time**.	The stream storage layer is responsible for providing **scalable** and **cost-effective** components to store streaming data. The streaming data can be **stored in the order it was received** for a **set duration of time** and can be **replayed indefinitely during that time**.	The stream processing layer is responsible for **transforming data into a consumable state** through data validation, cleanup, normalization, transformation, and enrichment. The streaming records are read in the order they are produced, allowing for real-time analytics, building event driven applications, or streaming ETL.	The destination layer is like a purpose-built destination depending upon your use case. Your destination can be an event driven application, data lake, data warehouse, database, or an OpenSearch.
• Smart meters • Database Change • Sensor networks • Social Media feeds	• HTTP/2 • WebSocket • MQTT (Message Queuing Telemetry Transport) • AMQP (Advanced Message Queuing Protocol)	• Kafka • RabbitMQ	• Kafka Streams • Spark • Flink • Storm	• Database (e.g MongoDB) • Storage (e.g E3) • Etc.

Source: *©Amazon - https://docs.aws.amazon.com/whitepapers/latest/build-modern-data-streaming-analytics-architectures/what-is-a-modern-streaming-data-architecture.html*

Main Requirements per Solution area

Technology Study

Stream Sources	Stream Ingestion	Stream Storage	Stream Processing	Stream Destination
	1. **Minimize complexity** for event producers to send their events by utilizing open standards and common interaction patterns; 2. Primarily support for up to **40.000.000 events per 15min**; 3. Authentication and authorization relying on **existing IAM implementation** (AD, certificates, partners); 4. **Minimize impact on existing platform** (from maintenance, development and deployment perspective), reuse existing components where possible.	1. Primarily support for up to **1.000.000.000 events per day** and technically scalable up to 5.000.000.000 events per day 2. **Open Source where possible** (no additional licenses); Third party (contracting)dependencies kept to the minimum to allow for flexible scaling, deployment and predictable costing. 3. **Stream to (micro) Batch capability** to integrate with MDM module; 4. **Minimize impact on existing platform** (from maintenance, development and deployment perspective), reuse existing components where possible.	1. **Stream to (micro) Batch capability** to integrate with MDM module; 2. **Minimize impact on existing platform** (from maintenance, development and deployment perspective), reuse existing components where possible. 3. **Open Source where possible** (no additional licenses); Third party (contracting)dependencies kept to the minimum to allow for flexible scaling, deployment and predictable costing.	

Technology Analysis

Technology Study

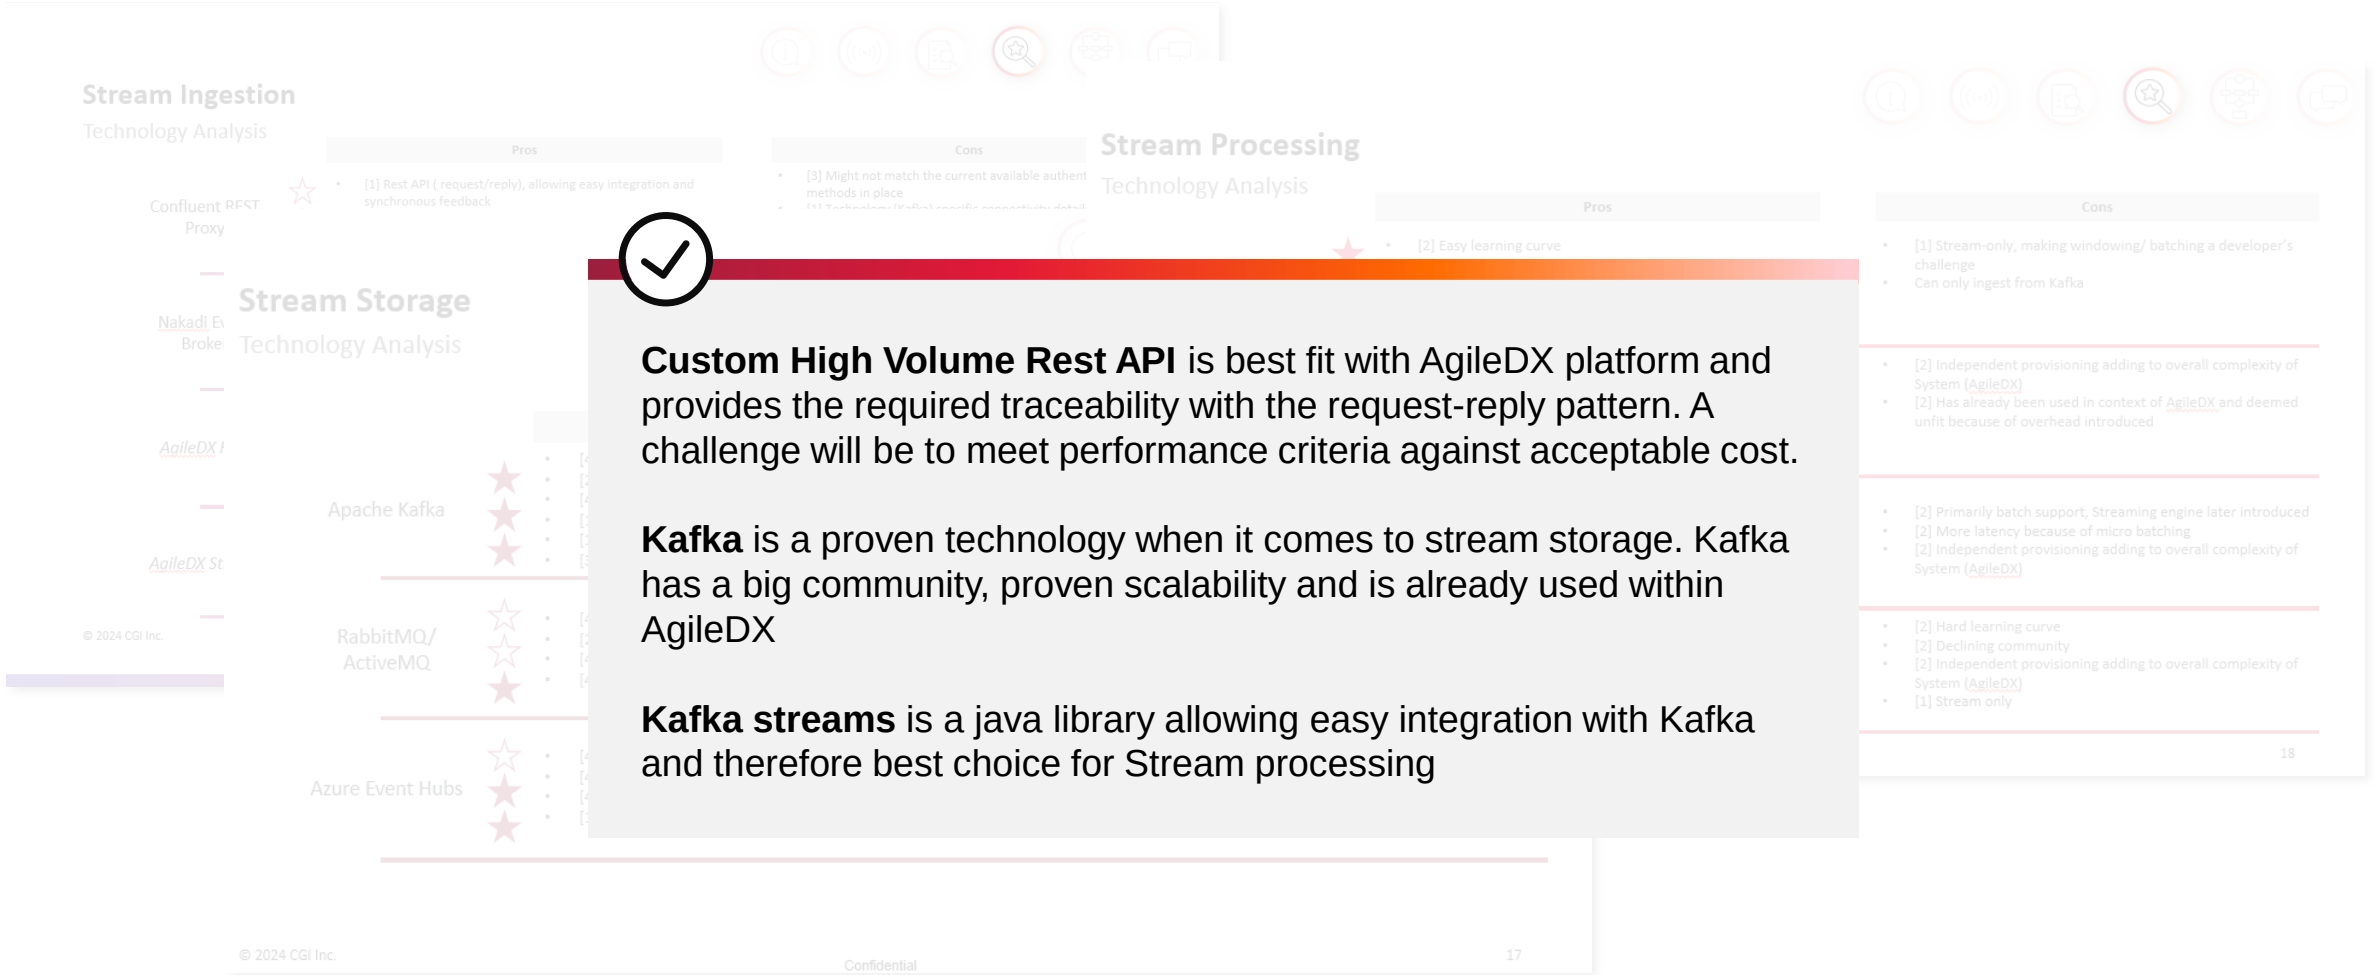

Custom High Volume Rest API is best fit with AgileDX platform and provides the required traceability with the request-reply pattern. A challenge will be to meet performance criteria against acceptable cost.

Kafka is a proven technology when it comes to stream storage. Kafka has a big community, proven scalability and is already used within AgileDX

Kafka streams is a java library allowing easy integration with Kafka and therefore best choice for Stream processing

High Level Solution Design

Technology Study

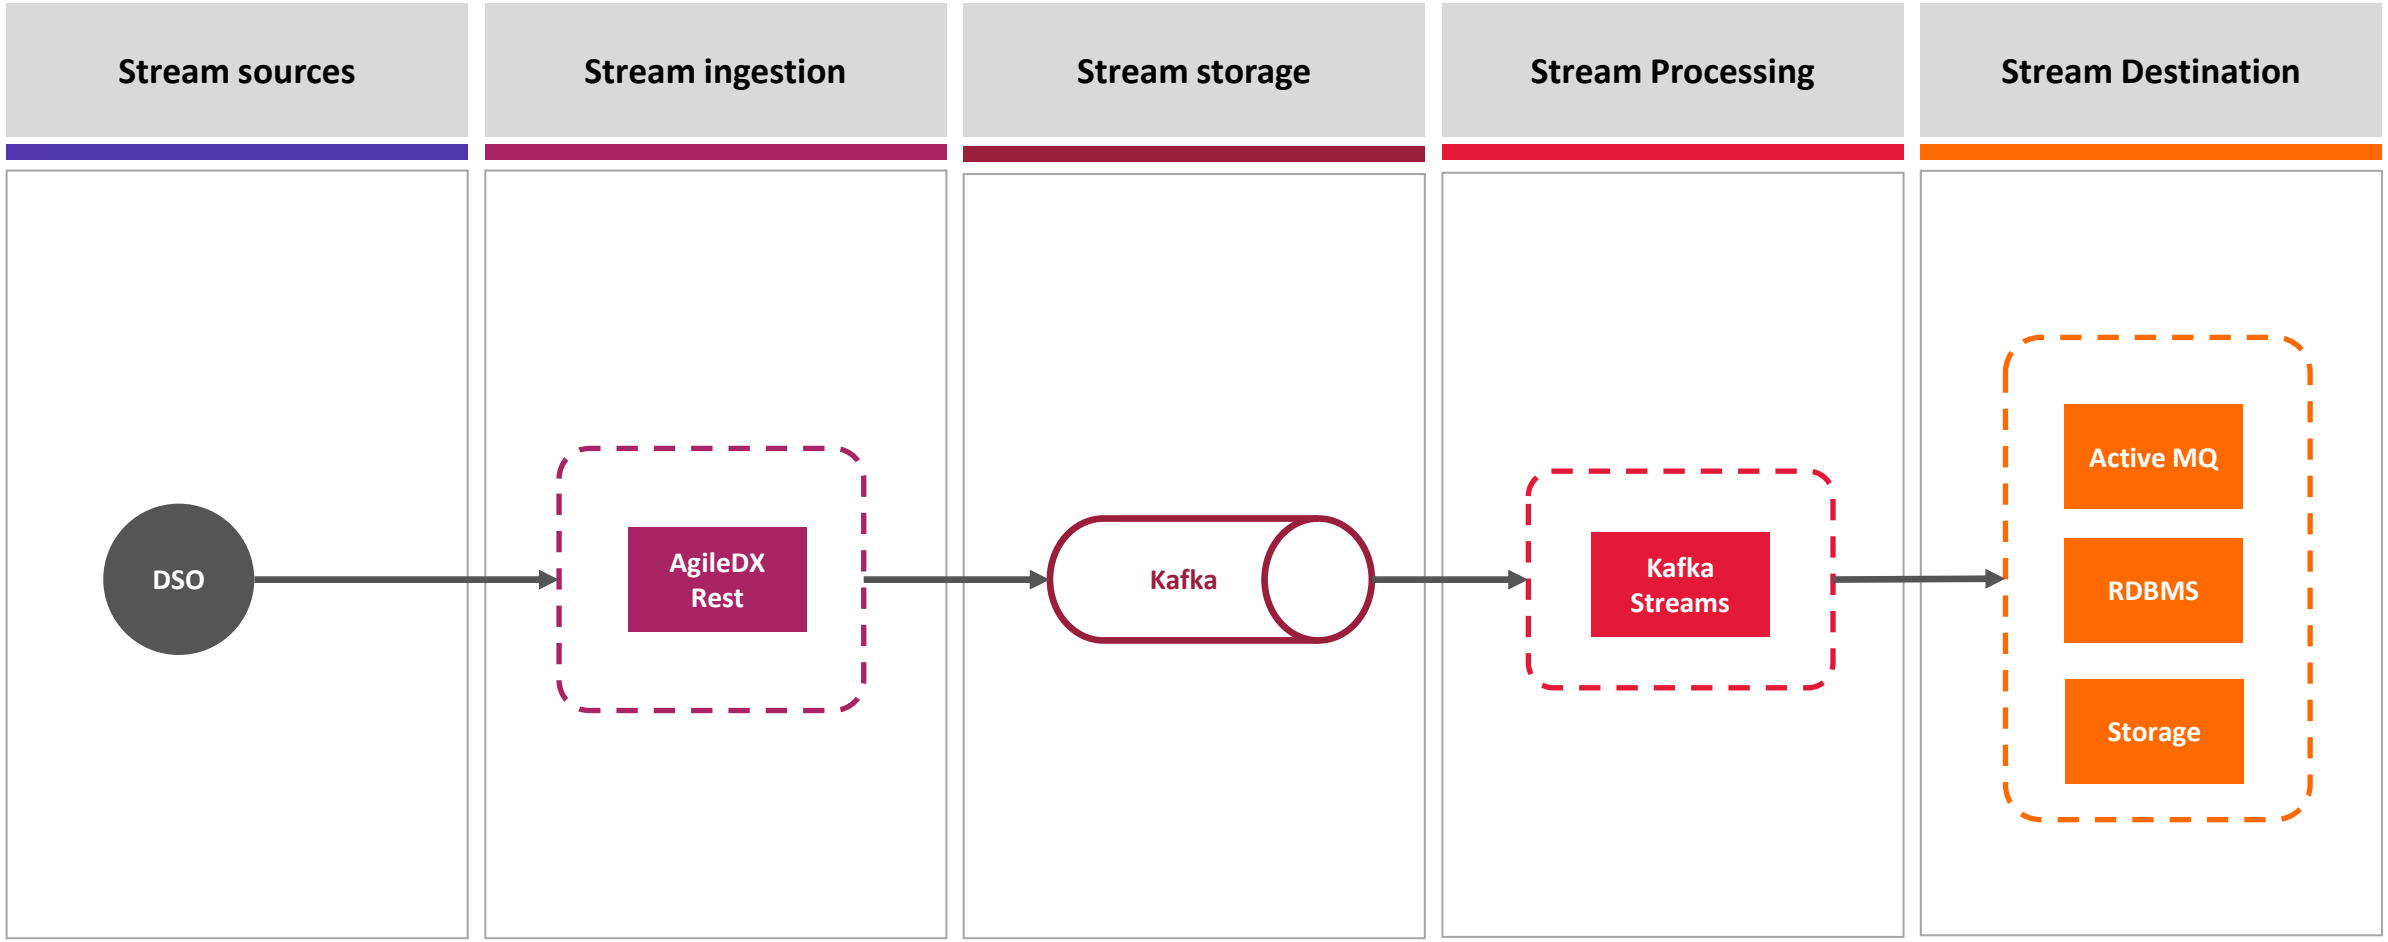

HTTP/2

Technology Study

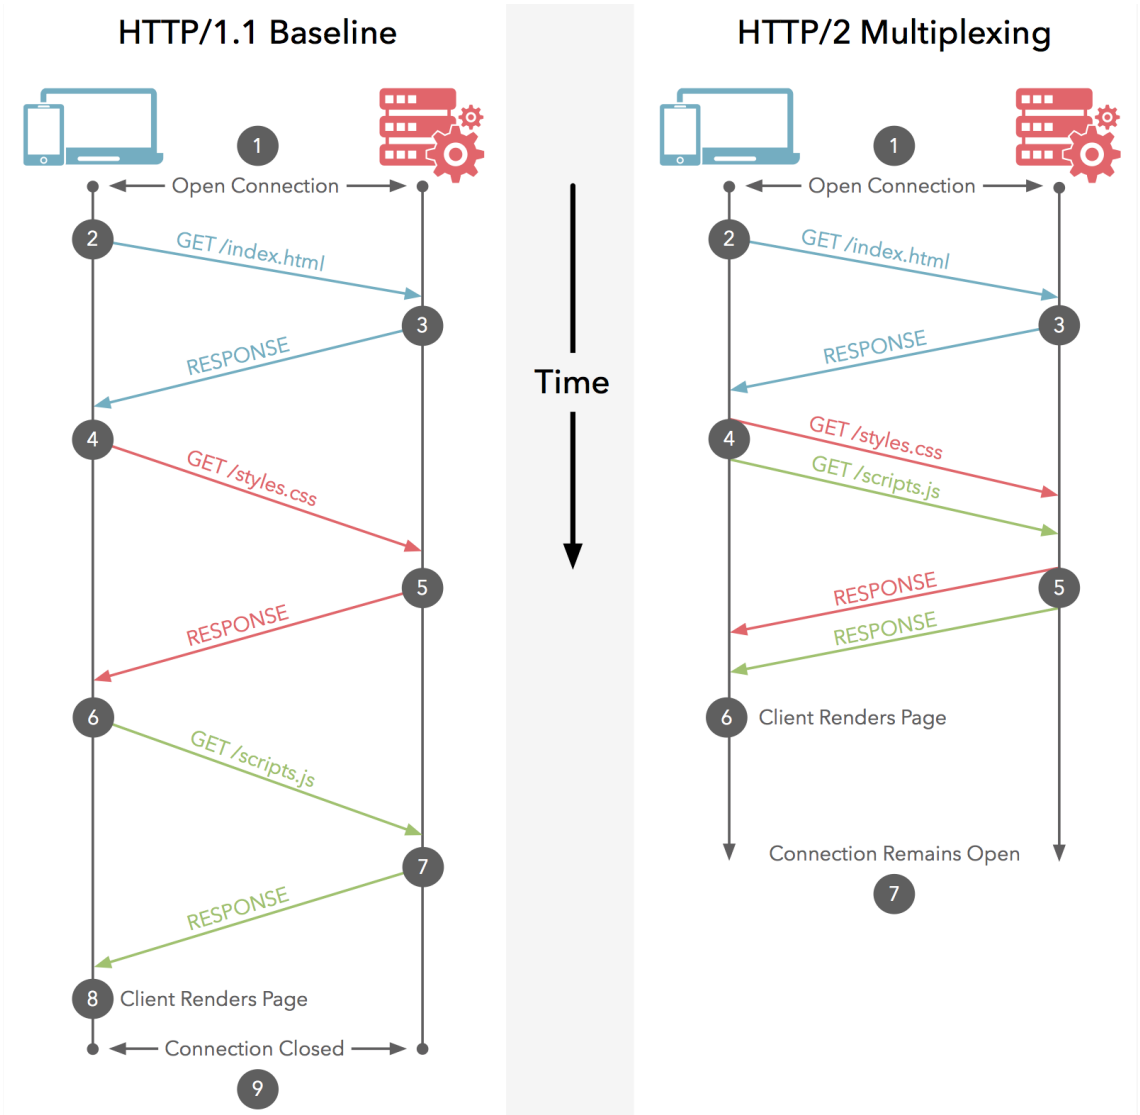

Proof of Concept

Phase 2

Proof of Concept

Stream ingestion

Proof of Concept

High performant ingestion channel integrated with available AgileDX logging and authentication modules. Verify pattern to be used in actual implementation

- [preferred] **REST API based on HTTP/2** preferred because of reduced complexity for participants and synchronous reply if a message has successfully been received.
- [Fallback] Stream API based on protocols like websocket, MQTT

Goal of the proof of concept is to confirm solution direction for receiving events with special attention to:

- Scalability (> 40 mio events per 15 minutes)
- Security integration
- Kafka architecture

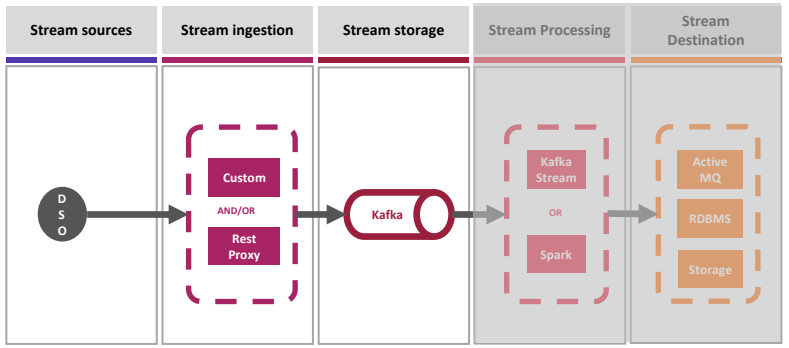

Stream Processing

Proof of Concept

Verify if **Kafka Streams** suffice as stream processor or additional technologies are required. If so, extend proof of concept with Apache Spark.

Goal of the proof of concept is to verify if Kafka Streams is able to process events and combine them to batch messages for MDM to process. Special attention goes to:

- Scalability (> 40 mio events per 15 minutes)
- Windowing capabilities
- Metadata requirements/constraints to allow for efficient Stream to Batch
- Baseline Kafka cluster configuration to allow Stream to Batch

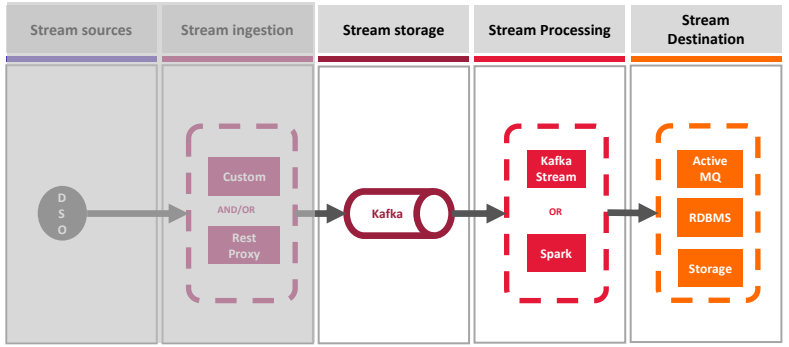

Work group

Introduction of a high volume event handling channel will introduce technical and possibly functional deviations from the current way of working. To allow optimal fit for the Finnish market it is advised to have a work group organized with experts from Fingrid, DSO's, Software vendors and CGI to agree on aspects like:

- Event format (attributes, identifiers)
- Patterns (Request/ Reply; Put and Update)
- Rejections and/or confirmations
- Distribution to other market parties.
- How to support specific scenarios like "energy communities"

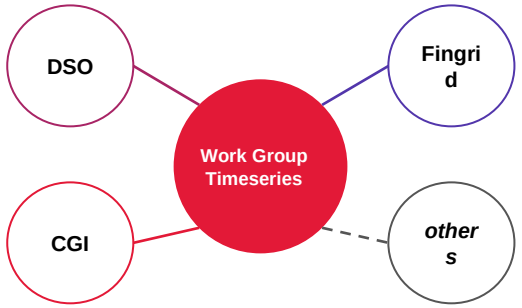

Planning

Proof of Concept

PoC Stream Ingestion	• Design/build/test/provision Stream ingestion • Loadtest Stream Ingestion • Optimize • Report
PoC Stream Processing	• Design/build/test/provision Stream processing • Loadtest Stream Processing • Optimize • Report
Workgroup	• Event format (attributes, identifiers) • Patterns (Request/ Reply; working model) • Rejections and/or confirmations • Distribution to other market parties. • How to support specific scenarios like “energy communities”

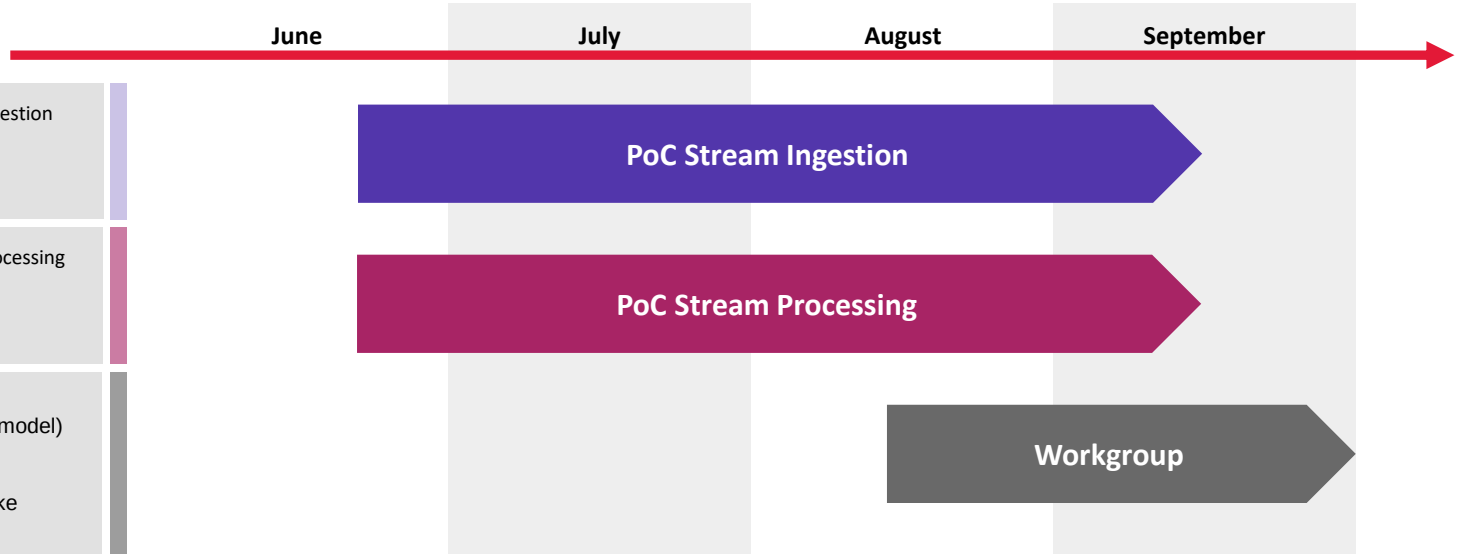

Workgroup Planning

Proof of Concept

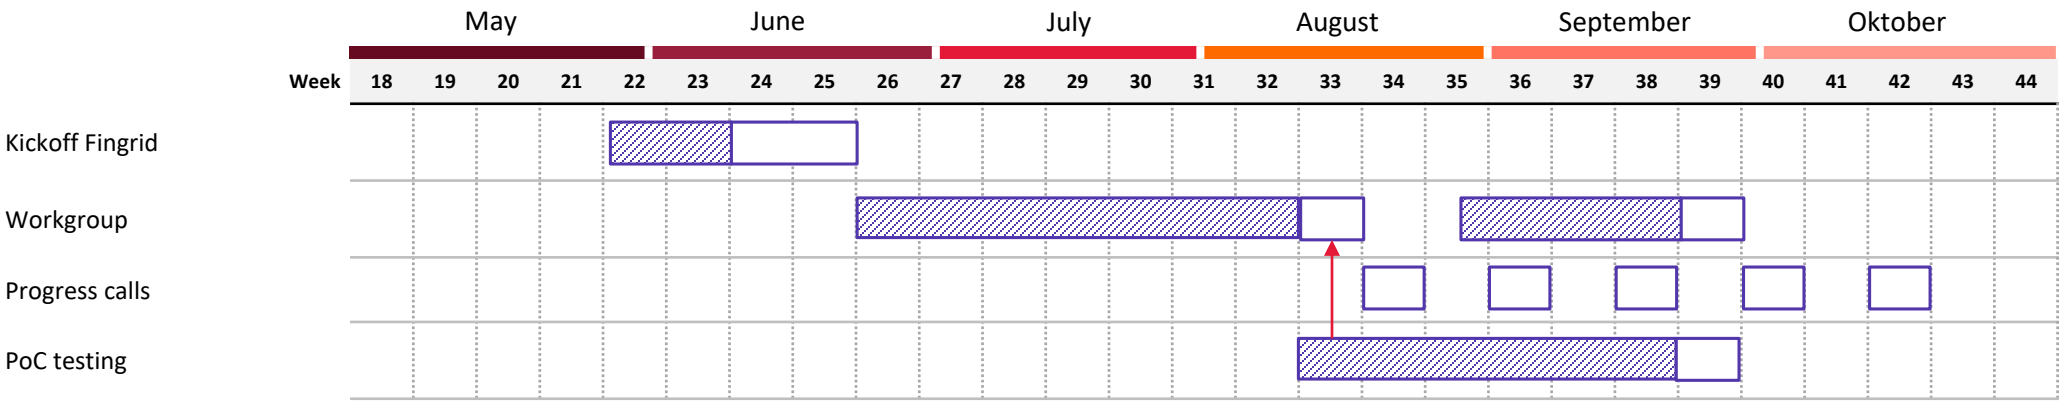

- Kick-off fingrid to align on previous steps (phase 1 – report), future steps (phase 2), Preparation for workshop
- Workshop with sector.
  - 1st (13-14 August): Kickoff, Event format, Pattern, Rejection/Confirmation, status PoC
  - 2nd (25th September): Lookback, Distribution to other market parties, Special cases, Testing days
- Progress calls: Biweekly digital call to address progress on actions from workshop (only when needed)
- PoC Testing: 2 days (23-24 September) onsite for connecting selected parties to sandbox. During 1st workshop the Test tooling as used by CGI will be made available to stakeholders.

Week meeting will be held

Lead/preparation time for participants. Supported with ad-hoc meetings if needed

Phase 2

Proof of Concept

PoC
CGI is responsible for the deliverable
Fingrid is informed

- **Report - Event Channel Timeseries** describing the results of the proof of concepts including the feasibility of the technology and a continuation plan in case technology is deemed unfeasible
- **High Level Design - Event Channel Timeseries** including design considerations, decisions and dependencies and expected on premise resource requirements, to be used as input for System Design/Architecture
- **Test Tool** (driver written in NodeJS) shared with workgroup, prior to testing days, as reference.

Workgroup
CGI (and Fingrid) responsible for the deliverable
Workgroup is consulted

- **Technical Interface Specification** where Fingrid and CGI are responsible for creating the draft and workgroup is consulted
- Memo per main topic addressed by workgroup
  - **Integration pattern** (Request/Response/ working model)
  - **Traceability** (confirmation/ rejection)
  - **Special cases** (e.g. Energy Communities)
  - **Distribution of Timeseries** to DSO and Suppliers
  - Concept **Test Planning** and environment dependencies phase 3

  These memos are to be used as input for System Design/Architecture